

TUC Transfer Protocol Daemon
&
TUC Alert Daemon:

System Documentation

Martin Thorsen Ranang

October 17, 2011

Contents

Contents	iii
Abstract	vii
1 Overview	1
1.1 TUC Transfer Protocol Daemon	1
1.2 TUC Alert Daemon	3
2 Features	5
2.1 Design Criteria	5
2.2 Scalability and Robustness	5
3 Program Usage	9
3.1 TUC Transfer Protocol Daemon	9
3.2 TUC Transfer Protocol Daemon Controller	9
3.3 TUC Transfer Protocol Client Application	11
A Module Application Program Interface	13
Bibliography	15

Abbreviations

API	application program interface
CPU	central processing unit
I/O	input/output
SMS	Short Message Service
TAD	TUC Alert Daemon
TTPD	TUC Transfer Protocol Daemon
TUC	The Understanding Computer
XML	Extensible Markup Language

Abstract

This document provides an overview of both the TUC Transfer Protocol Daemon and the TUC Alert Daemon software. The software package contains two programs, named `ttpd` and `ttpc`. These programs use a common application program interface, which is also documented herein.

Chapter 1

Overview

The main purpose of TUC Transfer Protocol Daemon (*TTPD*) is to work as a mediator for requests between external programs and The Understanding Computer (*TUC*). In addition, if used with the TUC Alert Daemon (*TAD*) it will also handle ‘alert’¹ requests when used in conjunction with the Short Message Service (*SMS*) interface. Figure 1.1 shows an overview of the system architecture.

Both programs were implemented in Python version 2.3.4.

1.1 TUC Transfer Protocol Daemon

The *TTPD* is implemented as a *threading* server, listening for connections from *clients* on a user-specified network port (by using a *socket*). When a request is received, the server starts a *handler thread* responsible for handling the request. First, the nature of the request is determined. If it is an ‘alert cancellation’ request, the request is handled by the handler thread through communication with *TAD* and the *client* (see Section 1.2 for more information about this behavior). However, if the request is not of that nature, it is forwarded to *TUC* for processing. This is done by having the server act as a *producer thread* that places incoming *tasks* in a *thread pool* with $n \in \langle 0, k \rangle$ *consumer threads*. Each of the consumer threads controls one external *TUC* process each.

The external *TUC* processes are started by the server before it starts accepting connections. This is done because it would take too long if an external *TUC* process should be started in order to handle each request. In relation

¹That is, messages like “Kan du varsle meg 15 minutter før bussen fra Nardo til Gløshaugen går?”

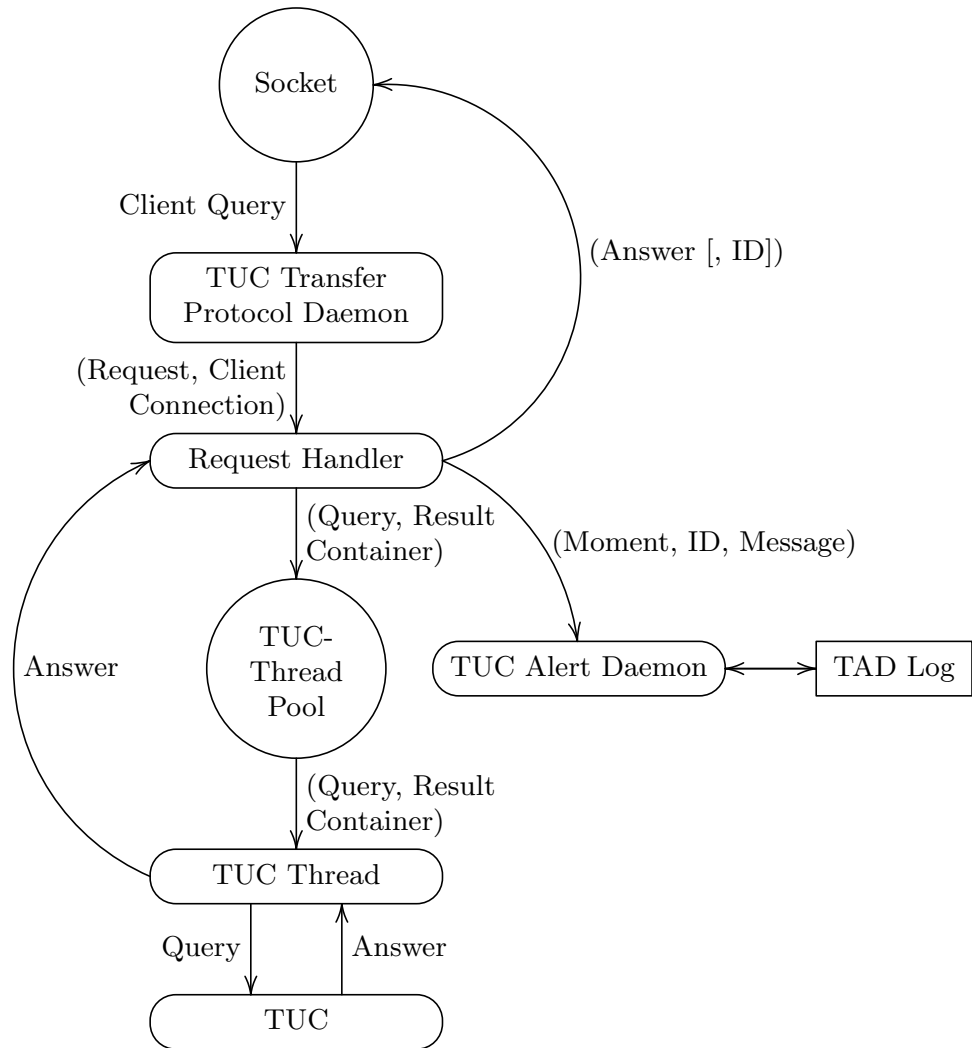


Figure 1.1: Overview of the TTPD and TAD system.

to the main server process, the TUC are *forked* processes. The communication between each *consumer* thread and its TUC process is done through *pipes* that function as the TUC process' `stdin` and `stdout` file streams.

When the result from TUC has been received, the *consumer thread* stores it in a *thread-safe* container that only the original *handler thread* and the consumer thread shares. Hence, the response for handling the request is again returned to the handler thread. Now, the result is parsed, the appropriate information is logged and an answer is given to the client.

1.2 TUC Alert Daemon

The TAD is built-in as a part of TTPD, but has its own responsibilities.

The most central component of TAD is a *scheduler* responsible for the timing of sending out alert-messages at the moments specified by users.

TAD consists mainly of two threads that for the most of the time run independently of the TTPD process. One thread constitutes an *alert scheduler* while the other thread handles requests of adding and removing alerts and is responsible for communicating with the *database engine*. The communication between a *TTPD request handler* and the *TAD request handler* is done through a shared request pool (thread-safely protected by *lock*).

Chapter 2

Features

This chapter presents the design criteria and features that characterizes TUC Transfer Protocol Daemon (*TTPD*) and TUC Alert Daemon (*TAD*).

2.1 Design Criteria

The main goals of this work was to make the system more robust and to implement an alert service. In addition to this, where there have been several possible solutions to a problem, the guiding principles has been as follows: choose solutions that are *easily maintainable* (use standard modules and programs if available, and write easily readable code and documentation), *highly scalable* and that will *probably not need to be changed* in the near future. In addition, the system has been designed in the spirit of *Unix* philosophy; embracing modularity, the use of stream redirection and process control as described in (Ranang 2003). For a very concise introduction to the theory and pragmatics of server/daemon design, please see (Griffin and Nelson 1998a; Griffin and Nelson 1998b; Griffin, Donnelly, and Nelson 1998).

2.2 Scalability and Robustness

First of all, the *scalability* of the system has been improved in several ways:

- The design of TTPD allows a computer to serve multiple The Understanding Computer (*TUC*) requests in parallel. This is ensured through the use of the pool of *TUC-threads* where each thread controls an externally running TUC process. This means that if more central processing unit (*CPU*) are added to the computer, each processes

can run on its own CPU. Even without multiple CPU, the computer can serve multiple requests in multiplexed parallel (but possibly with a longer delay than when serving single requests, depending both on CPU- and input/output (I/O) intensity).

- By combining the use of *socket* with *tread* (could have used *forked processes* but they take longer to create) the *daemon* is able to stack up incoming requests and serve them as soon as possible. This feature ensures that the clients can deliver their requests even when other requests are already being processed.

Secondly, the following features have been implemented to ensure a high level of *robustness* of the system:

- The use of *socket* ensure that *if* the daemon process should crash, any transaction received but not finished will silently be lost. That is, the log will contain information about the reception of the requests, but when the daemon is restarted, it will not result in a bombardment of the users with their old requests. An exception from this is the handling of any due TAD alerts. They will be sent out as soon as possible after the alert moment have passed.
- Efforts have been made to graciously handle extremely high workloads. The daemon has a configurable “soft” high-load limit. This results in a predefined message being sent back to users—when the current number of concurrently handled requests is above this limit—telling the user that due to very high demand, the request cannot be handled at this moment.
- The daemon tries to cope with non-perfect conditions in a controlled way:
 - If one of the encapsulated TUC processes die unexpectedly, the daemon will start a replacement process immediately (this constitutes a built-in *watch dog* feature).
 - All alerts that are registered by the TAD is immediately saved to disk. This is done to ensure that during startup, the daemon consults the old log file to restore its state.
 - Because the *protocol* for *communication* with the Short Message Service (SMS) *message switch* defines the communication as asynchronous, it is possible that the daemon experiences situations when it cannot connect to the remote server to deliver its answer. If such a situation should arise—and it has, during intensive testing of the system—a *retry/resend algorithm* has been implemented. The *handler* of that request will sleep for a random number of seconds (within a predefined interval) and

then try to resend the answer when it awakens. If it still does not succeed, it repeats this behavior until a total time limit has been reached (suggestively sat to 15 minutes).

- Since the *SMS message switch communication protocol* uses Extensible Markup Language (*XML*)-based headers, the daemon parses these with the default *Python SAX-parser* (usually an *expat parser*). This ensures a robust handling of incoming messages and an automatic check for XML compliance.
- The TTPD is designed to handle random connections to its socket, so that non-protocol clients will not crash it.

Chapter 3

Program Usage

The programs `ttpd` (the *daemon*) and `ttpc` (the client interface application) both contain some built-in help information that is always available from the command line. All you need to do is to supply one of the flags `-help` or `-h` to the program on the command line.

3.1 TUC Transfer Protocol Daemon

The help menu in `ttpd` looks like:

3.2 TUC Transfer Protocol Daemon Controller

The help menu in `ttpdctl` looks like:

Usage: `ttpdctl [options] start|stop|restart|store_alerts`

Options:

<code>--version</code>	show program's version number and exit
<code>-h, --help</code>	show this help message and exit
<code>-c FILENAME, --config-file=FILENAME</code>	the FILENAME containing the default configuration; the default is <code>'/etc/default/ttpd'</code>
<code>-a ADDRESS, --ip-address=ADDRESS</code>	the ADDRESS of the interface on which to listen for connections from clients; if ADDRESS is 0 (the default) the server will listen on all available network interfaces
<code>-p PORT, --port=PORT</code>	listen for connections from clients on PORT; the

```

                                default is 2004
-A ADDRESS, --remote-ip-address=ADDRESS
                                the ADDRESS of the interface on which to send outgoing
                                communication; default is 'localhost'
-P PORT, --remote-port=PORT
                                the remote PORT; the default is 2005
--service-name=SMS_SERVICE_NAME
                                the SERVICE_NAME for the SMS service. The default is
                                'ROUTE'.
-q SIZE, --socket-queue-size=SIZE
                                the SIZE of the socket queue; the default is 5
-Q THREADS, --high-load-limit=THREADS
                                maximum number of concurrently running THREADS; the
                                default is 50
-L LEVEL, --log-level=LEVEL
                                the filter LEVEL used when logging; possible values
                                are 'notset' < 'protocol' < 'debug' < 'info' <
                                'warning' < 'error' < 'critical'; the default is
                                'info'
-f FILE, --log-file=FILE
                                store the log in FILE; the default is 'ttpdctl.log'
-t, --without-tad
                                don't start the TUC Alert Daemon (TAD)
-T, --without-tuc
                                don't start the TUC processes
-n NUMBER, --tuc-threads=NUMBER
                                start NUMBER concurrent TUC processes; the default is
                                3
-s COMMAND, --path-to-tuc=COMMAND
                                run COMMAND as TUC subprocess; the default is
                                ./busestuc.sav
-E ENCODING, --tuc-external-encoding=ENCODING
                                tell SICStus Prolog that it runs in an ENCODING
                                environment. If "None", TTPD tries to set it
                                according to its stdin encoding. The default is None
-d, --no-daemon
                                don't run as a daemon process
--pid-file=FILENAME
                                store the pid of the server process in FILENAME; the
                                default is /var/run/ttpd/ttpd.pid
-g PROVIDER, --sms-gateway=PROVIDER
                                the PROVIDER to use when sending SMS; possible values
                                are 'esolutions', 'payex' and 'pswincom'. The default
                                is 'payex'.
-U USERNAME, --user=USERNAME
                                the USERNAME that the daemon process will run as. The
                                default is 'ttpd'.
-G GROUPNAME, --group=GROUPNAME
                                the GROUPNAME that the daemon process will run as.
                                The default is 'ttpd'.
--originating-address=PHONE_NUMBER
                                the PHONE_NUMBER of the SMS service. The default is
                                '1939'.

```

```

--payex-test          use the PayEx test, instead of the production,
                      interface. The default is 'False'.
--payex-trace=FILENAME
                      store PayEx communication traces in FILENAME. If
                      option is not used, no trace will be performed.
--payex-key=ENCRYPTION_KEY
                      the ENCRYPTION_KEY to use when sending SMS via PayEx
                      interface. The default is 'PgHzip4b2RH8u43XSE6V'.
--payex-account=ACCOUNT_NUMBER
                      the ACCOUNT_NUMBER to use when sending SMS via PayEx
                      interface. The default is '21217859'.
--payex-test-key=ENCRYPTION_KEY
                      the ENCRYPTION_KEY to use for sending SMS via PayEx
                      test interface. The default is
                      '2PuM2YTbK3VfUypbN2bU'.
--payex-test-account=ACCOUNT_NUMBER
                      the ACCOUNT_NUMBER to use when sending SMS via PayEx
                      test interface. The default is '50017893'.
-o FILE, --old-log-file=FILE
                      location of the old log FILE, used to retrieve the PID
                      of the current process, when restarting after a file
                      has been moved (e.g., log rotated)
-e EXECUTABLE, --executable=EXECUTABLE
                      location of the ttpd EXECUTABLE; the default is ttpd

```

3.3 TUC Transfer Protocol Client Application

The help menu in `ttpc` looks like:

Usage: `ttpc [options] [request]`

Options:

```

--version            show program's version number and exit
-h, --help          show this help message and exit
-c FILENAME, --config-file=FILENAME
                    the FILENAME containing the default configuration; the
                    default is '/etc/default/ttpd'
-a ADDRESS, --ip-address=ADDRESS
                    the ADDRESS of the interface on which to listen for
                    connections from clients; if ADDRESS is 0 (the
                    default) the server will listen on all available
                    network interfaces
-p PORT, --port=PORT listen for connections from clients on PORT; the
                    default is 2004
-A ADDRESS, --remote-ip-address=ADDRESS
                    the ADDRESS of the interface on which to send outgoing
                    communication; default is 'localhost'
-P PORT, --remote-port=PORT

```

```

        the remote PORT; the default is 2005
--service-name=SMS_SERVICE_NAME
        the SERVICE_NAME for the SMS service. The default is
        'ROUTE'.
-q SIZE, --socket-queue-size=SIZE
        the SIZE of the socket queue; the default is 5
-Q THREADS, --high-load-limit=THREADS
        maximum number of concurrently running THREADS; the
        default is 50
-L LEVEL, --log-level=LEVEL
        the filter LEVEL used when logging; possible values
        are 'notset' < 'protocol' < 'debug' < 'info' <
        'warning' < 'error' < 'critical'; the default is
        'info'
-f FILE, --log-file=FILE
        store the log in FILE; the default is 'ttpc.log'
-I TRANS_ID, --transaction-id=TRANS_ID
        the TRANS_ID to use when communicating with a remote
        server. The default is 'LINGSMSIN'.
-l, --listen-mode
        listen for inbound connections
-t, --test
        run in looping test-mode
-T PHONE, --phone-number=PHONE
        the number that the message supposedly was sent from
        the default is 'None'
-F, --fake-outgoing
        fake an outgoing message from the server
-i FILE, --input-file=FILE
        read input from FILE
-w, --web
        generate output suitable for the Web; this implicates
        --transaction-id=WEB
-j, --json
        generate JSON output suitable for Web Services; this
        implicates --transaction-id=JSON
-W, --show-technical
        show technical information (semantics)

```

Appendix A

Module Application Program Interface

It should be noted that the SocketServer module was not developed by the author, but since the modules TTP.Server and TTP.Handler contain classes that inherit from classes in SocketServer, it has been included here to make the application program interface (*API*) documentation meaningful.

Bibliography

- [Griffin, Donnelly, and Nelson1998] Griffin, Ivan, Mark Donnelly, and John Nelson.
1998.
“Linux Network Programming, Part 3.”
Linux J. 1998 (48es): 7 (April).
- [Griffin and Nelson1998a] Griffin, Ivan, and John Nelson.
1998a.
“Linux Network Programming, Part 1.”
Linux J. 1998 (46es): 5 (February).
- [Griffin and Nelson1998b] ———.
1998b.
“Linux Network Programming, Part 2.”
Linux J. 1998 (47es): 6 (March).
- [Ranang2003] Ranang, Martin Thorsen.
2003.
“The Fragrance of Unix.”
Compendium.
Version 1.1, available on-line.

